

# Homework 2

Pedro Sousa Meireles

October 22, 2018

```
In [1]: import numpy
        from urllib.request import urlopen
        import scipy.optimize
        import random
        import ast
        from sklearn.decomposition import PCA
        from collections import defaultdict
        import networkx as nx
        import matplotlib.pyplot as plt
        from sklearn.metrics import precision_recall_curve
        from sklearn.utils.fixes import signature
        from math import exp
        from math import log
```

## 0.0.1 Question 1

```
In [2]: def parseData(fname):
        for l in open(fname):
            yield ast.literal_eval(l)
```

```
In [3]: print("Reading data...")
        data = list(parseData("beer_50000.json"))
        print("done")
```

Reading data...  
done

```
In [4]: def feature(datum):
        feat = [1, datum['review/taste'], datum['review/appearance'], datum['review/aroma'],
                datum['review/palate'], datum['review/overall']]
        return feat
```

```
X = [feature(d) for d in data]
y = [d['beer/ABV'] >= 6.5 for d in data]
```

```
In [5]: def inner(x,y):
        return sum([x[i]*y[i] for i in range(len(x))])
```

```

def sigmoid(x):
    return 1.0 / (1 + exp(-x))

In [6]: #####
# Logistic regression by gradient ascent #
#####

# NEGATIVE Log-likelihood
def f(theta, X, y, lam):
    loglikelihood = 0
    for i in range(len(X)):
        logit = inner(X[i], theta)
        loglikelihood -= log(1 + exp(-logit))
        if not y[i]:
            loglikelihood -= logit
    for k in range(len(theta)):
        loglikelihood -= lam * theta[k]*theta[k]
    # for debugging
    # print("ll =" + str(loglikelihood))
    return -loglikelihood

# NEGATIVE Derivative of log-likelihood
def fprime(theta, X, y, lam):
    dl = [0]*len(theta)
    for i in range(len(X)):
        logit = inner(X[i], theta)
        for k in range(len(theta)):
            dl[k] += X[i][k] * (1 - sigmoid(logit))
            if not y[i]:
                dl[k] -= X[i][k]
        for k in range(len(theta)):
            dl[k] -= lam*2*theta[k]
    return numpy.array([-x for x in dl])

ind_list = [i for i in range(len(X))]
random.shuffle(ind_list)
random_X = []
random_y = []

for i in ind_list:
    random_X.append(X[i])
    random_y.append(y[i])

X_train = random_X[:len(random_X)//3]
y_train = random_y[:len(random_y)//3]

X_test = random_X[len(random_X)//3 + 1:2*len(random_X)//3]

```

```

y_test = random_y[len(random_y)//3 + 1:2*len(random_y)//3]

X_val = random_X[2*len(random_X)//3:]
y_val = random_y[2*len(random_y)//3:]

In [7]: #####
# Train #
#####

def train(lam):
    theta,_,_ = scipy.optimize.fmin_l_bfgs_b(f, [0]*len(X[0]), fprime, pgtol = 10, args =
    return theta

In [8]: #####
# Predict #
#####

def performance(theta, X_test, y_test):
    scores = [inner(theta,x) for x in X_test]
    predictions = [s > 0 for s in scores]
    correct = [(a==b) for (a,b) in zip(predictions,y_test)]
    acc = sum(correct) * 1.0 / len(correct)
    return predictions, acc, scores

In [9]: #####
# Validation pipeline #
#####

lam = 1.0

theta = train(lam)
train_predictions, train_acc, train_scores = performance(theta, X_train, y_train)
test_predictions, test_acc, test_scores = performance(theta, X_test, y_test)
val_predictions, val_acc, val_scores = performance(theta, X_val, y_val)
print("lambda = " + str(lam))
print("Train:\taccuracy=" + str(train_acc))
print("Test:\taccuracy=" + str(test_acc))
print("Validation:\taccuracy=" + str(val_acc))

lambda = 1.0
Train:      accuracy=0.7214088563542542
Test:       accuracy=0.7146285851434058
Validation: accuracy=0.7153056938861223

```

## 0.0.2 Question 2

```

In [16]: tPos = 0
         fPos = 0

```

```

tNeg = 0
fNeg = 0

for i in range(len(y_test)):
    if y_test[i] == True and test_predictions[i] == True:
        tPos += 1
    elif y_test[i] == False and test_predictions[i] == True:
        fPos += 1
    elif y_test[i] == False and test_predictions[i] == False:
        tNeg += 1
    else:
        fNeg += 1

print("True Positive: " + str(tPos))
print("False Positive: " + str(fPos))
print("True Negative: " + str(tNeg))
print("False Negative: " + str(fNeg))

```

```

True Positive: 9019
False Positive: 3504
True Negative: 2891
False Negative: 1252

```

### 0.0.3 Question 3

```

In [17]: test_predictions = list(zip(y_test, test_predictions, test_scores))
        test_predictions = sorted(test_predictions, key=lambda x:x[2])

In [18]: relevantAt100 = sum([1 for v in test_predictions[-100:] if v[0] == True])
        retrievedAt100 = sum([1 for v in test_predictions[-100:] if v[1] == True])
        relRetrievedAt100 = sum([1 for v in test_predictions[-100:]
                                if v[0] == v[1] and v[0] == True])
        totalRelevant = sum([1 for v in test_predictions if v[0] == True])
        totalRetrieved = sum([1 for v in test_predictions if v[1] == True])
        totalRelRetrieved = sum([1 for v in test_predictions
                                if v[0] == v[1] and v[0] == True])

In [19]: precisionAt100 = relRetrievedAt100/retrievedAt100
        recallAt100 = relRetrievedAt100/relevantAt100
        precision = totalRelRetrieved/totalRetrieved
        recall = totalRelRetrieved/totalRelevant

In [20]: print("Precision@100: " + str(precisionAt100))
        print("Precision: " + str(precision))
        print("Recall@100: " + str(recallAt100))
        print("Recall: " + str(recall))

Precision@100: 0.92
Precision: 0.7201948414916554

```


```
Recall@100: 1.0  
Recall: 0.8781033979164639
```

#### 0.04 Question 4

```
In [21]: precision, recall, _ = precision_recall_curve(y_test, test_scores)
```

```
In [22]: plt.step(recall, precision, color='b', alpha=0.2,  
               where='post')  
         plt.fill_between(recall, precision, alpha=0.2, color='b')  
  
         plt.xlabel('Recall')  
         plt.ylabel('Precision')  
         plt.ylim([0.0, 1.05])  
         plt.xlim([0.0, 1.0])  
         plt.title('2-class Precision-Recall curve: AP={0:0.2f}')
```

```
Out[22]: Text(0.5,1,'2-class Precision-Recall curve: AP={0:0.2f}')
```



#### 0.05 Question 5

```
In [23]: ### Network visualization ###
```

```

edges = set()
nodes = set()
for edge in open("egonet.txt"):
    x,y = edge.split()
    x,y = int(x),int(y)
    edges.add((x,y))
    edges.add((y,x))
    nodes.add(x)
    nodes.add(y)

G = nx.Graph()
for e in edges:
    G.add_edge(e[0],e[1])
nx.draw(G)
plt.show()
plt.clf()

```



<Figure size 432x288 with 0 Axes>

```
In [24]: connComponents = list(nx.connected_component_subgraphs(G))
```

```
In [25]: for graph in connComponents:
          print(len(graph.nodes))
```

40  
9  
12

The graph has 3 connected components and the largest one has 40 nodes

### 0.0.6 Question 6

```
In [26]: largestCom = connComponents[0]

In [27]: nodes = list(largestCom.nodes)
         edges = list(largestCom.edges)

In [28]: nodes1 = nodes[:len(nodes)//2]
         degrees1 = sum([G.degree(v) for v in nodes1])
         nodes2 = nodes[len(nodes)//2:]
         degrees2 = sum([G.degree(v) for v in nodes2])

In [29]: cuttedEdges = sum([1 for e in edges
if (e[0] in nodes1 and e[1] in nodes2) or (e[0] in nodes2 and e[1] in nodes1)])

In [30]: normCut = cuttedEdges/2*(1/degrees1 + 1/degrees2)

In [31]: print("Normalized-cut cost: " + str(normCut))

Normalized-cut cost: 0.39416666666666667
```

### 0.0.7 Question 7

```
In [32]: newCost = 0
         while 1:
             costs = []
             for n in nodes:
                 newNodes1 = nodes1.copy()
                 newNodes2 = nodes2.copy()

                 if n in nodes1:
                     newNodes1.remove(n)
                     newNodes2.append(n)
                     newDegrees1 = degrees1 - G.degree(n)
                     newDegrees2 = degrees2 + G.degree(n)
                 else:
                     newNodes2.remove(n)
                     newNodes1.append(n)
                     newDegrees2 = degrees2 - G.degree(n)
                     newDegrees1 = degrees1 + G.degree(n)
```

```

        cuttedEdges = sum([1 for e in edges
                           if (e[0] in newNodes1 and e[1] in newNodes2)
                           or (e[0] in newNodes2 and e[1] in newNodes1)])
        costs.append([n, cuttedEdges/2*(1/newDegrees1 + 1/newDegrees2)])

costs = sorted(costs, key=lambda x:x[1])
newCost = costs[0][1]
if costs[0][0] in nodes1:
    nodes1.remove(costs[0][0])
    nodes2.append(costs[0][0])
    degrees1 = degrees1 - G.degree(costs[0][0])
    degrees2 = degrees2 + G.degree(costs[0][0])
else:
    nodes2.remove(costs[0][0])
    nodes1.append(costs[0][0])
    degrees2 = degrees2 - G.degree(costs[0][0])
    degrees1 = degrees1 + G.degree(costs[0][0])

nodes1.sort()
nodes2.sort()

if newCost > normCut:
    break
else:
    normCut = newCost

```

In [33]: `print(nodes1)`

```
[697, 703, 708, 713, 719, 745, 747, 753, 769, 772, 774, 800, 803, 805, 810,
811, 819, 823, 828, 830, 840, 856, 869, 880, 890]
```

In [34]: `print(nodes2)`

```
[729, 798, 804, 825, 861, 863, 864, 876, 878, 882, 884, 886, 888, 889, 893]
```

In [35]: `print("The new normalized cut cost is " + str(normCut))`

The new normalized cut cost is 0.09817045961624274

## 0.0.8 Question 8

```

In [36]: nodes1 = nodes[:len(nodes)//2]
        degrees1 = sum([G.degree(v) for v in nodes1])
        nodes2 = nodes[len(nodes)//2:]
        degrees2 = sum([G.degree(v) for v in nodes2])

```



```

In [37]: def modularity(c1, c2):
    e00 = sum([1 for e in edges if e[0] in c1 and e[1] in c1])/len(edges)
    e11 = sum([1 for e in edges if e[0] in c2 and e[1] in c2])/len(edges)

    a0 = 0
    a1 = 0
    for e in edges:
        if(e[0] in c1):
            a0 = a0 + 1
        else:
            a1 = a1 + 1

        if(e[1] in c1):
            a0 = a0 + 1
        else:
            a1 = a1 + 1
    a0 = a0/(2*len(edges))
    a1 = a1/(2*len(edges))

    return e00 - a0*a0 + e11 - a1*a1

In [38]: actualModularity = modularity(nodes1, nodes2)
    newModularity = 1
    while 1:
        modularities = []
        for n in nodes:
            newNodes1 = nodes1.copy()
            newNodes2 = nodes2.copy()

            if n in nodes1:
                newNodes1.remove(n)
                newNodes2.append(n)
                newDegrees1 = degrees1 - G.degree(n)
                newDegrees2 = degrees2 + G.degree(n)
            else:
                newNodes2.remove(n)
                newNodes1.append(n)
                newDegrees2 = degrees2 - G.degree(n)
                newDegrees1 = degrees1 + G.degree(n)

            modularities.append([n, modularity(newNodes1, newNodes2)])

        modularities = sorted(modularities, key=lambda x:x[1])
        newModularity = modularities[0][1]

    if modularities[0][0] in nodes1:
        nodes1.remove(modularities[0][0])
        nodes2.append(modularities[0][0])

```

```

        degrees1 = degrees1 - G.degree(modularities[0][0])
        degrees2 = degrees2 + G.degree(modularities[0][0])
    else:
        nodes2.remove(modularities[0][0])
        nodes1.append(modularities[0][0])
        degrees2 = degrees2 - G.degree(modularities[0][0])
        degrees1 = degrees1 + G.degree(modularities[0][0])

    nodes1.sort()
    nodes2.sort()

    if newModularity > actualModularity:
        break
    else:
        actualModularity = newModularity

```

In [39]: `print(nodes1)`

```
[697, 703, 708, 729, 745, 769, 772, 774, 803, 804, 810, 811, 819, 823, 825, 830, 863, 864,
876, 889, 893]
```

In [40]: `print(nodes2)`

```
[713, 719, 747, 753, 798, 800, 805, 828, 840, 856, 861, 869, 878, 880, 882, 884, 886, 888, 890]
```

In [41]: `print("The new modularity is " + str(actualModularity))`

The new modularity is -0.11372933884297518